



Bias Correction + XGBoost

by [Tili](#) · last run 2 days ago · Python script · 2019 views
using data from [Allstate Claims Severity](#) · Public

37

voters

Code

This script has been released under the [Apache 2.0](#) open source license.[Download Code](#)

```
1  import numpy as np
2  import pandas as pd
3  from datetime import datetime
4  from sklearn.preprocessing import StandardScaler
5  from sklearn.cross_validation import KFold
6  from sklearn.metrics import mean_absolute_error
7  from scipy.stats import skew, boxcox
8  from math import exp, log
9  import xgboost as xgb
10
11
12  def timer(start_time=None):
13      if not start_time:
14          start_time = datetime.now()
15          return start_time
16      elif start_time:
17          tmin, tsec = divmod((datetime.now() - start_time).total_seconds(),
18 60)
19          print(' Time taken: %i minutes and %s seconds.' %
20                (tmin, round(tsec, 2)))
21
22
23  def scale_data(X, scaler=None):
24      if not scaler:
25          scaler = StandardScaler()
26          scaler.fit(X)
27      X = scaler.transform(X)
28      return X, scaler
29
30
```

```

31 DATA_TRAIN_PATH = '../input/train.csv'
32 DATA_TEST_PATH = '../input/test.csv'
33
34
35 def load_data(path_train=DATA_TRAIN_PATH, path_test=DATA_TEST_PATH):
36     train_loader = pd.read_csv(path_train, dtype={'id': np.int32})
37     train = train_loader.drop(['id', 'loss'], axis=1)
38     test_loader = pd.read_csv(path_test, dtype={'id': np.int32})
39     test = test_loader.drop(['id'], axis=1)
40     ntrain = train.shape[0]
41     ntest = test.shape[0]
42     train_test = pd.concat((train, test)).reset_index(drop=True)
43     numeric_feats = train_test.dtypes[train_test.dtypes != "object"].index
44
45     # compute skew and do Box-Cox transformation
46     skewed_feats = train[numeric_feats].apply(lambda x: skew(x.dropna()))
47     print("\nSkew in numeric features:")
48     print(skewed_feats)
49     # transform features with skew > 0.25 (this can be varied to find
50 optimal value)
51     skewed_feats = skewed_feats[skewed_feats > 0.25]
52     skewed_feats = skewed_feats.index
53     for feats in skewed_feats:
54         train_test[feats] = train_test[feats] + 1
55         train_test[feats], lam = boxcox(train_test[feats])
56     features = train.columns
57     cats = [feat for feat in features if 'cat' in feat]
58     # factorize categorical features
59     for feat in cats:
60         train_test[feat] = pd.factorize(train_test[feat], sort=True)[0]
61     x_train = train_test.iloc[:ntrain, :]

```

Code

Output (3)

Comments (21)

Log

Versions (1)

Forks (13)

Fork Script

```

63 train_test_scaled, scaler = scale_data(train_test,
64 train, _ = scale_data(x_train, scaler)
65 test, _ = scale_data(x_test, scaler)
66
67 train_labels = np.log(np.array(train_loader['loss']))
68 train_ids = train_loader['id'].values.astype(np.int32)
69 test_ids = test_loader['id'].values.astype(np.int32)
70
71 return train, train_labels, test, train_ids, test_ids
72
73 ##### Actual Run Code
74 #####
75
76 # enter the number of folds from xgb.cv
77 folds = 5

```

```

78 cv_sum = 0
79 early_stopping = 25
80 fpred = []
81 xgb_rounds = []
82
83 start_time = timer(None)
84
85 # Load data set and target values
86 train, target, test, _, ids = load_data()
87 d_train_full = xgb.DMatrix(train, label=target)
88 d_test = xgb.DMatrix(test)
89
90 # set up KFold that matches xgb.cv number of folds
91 kf = KFold(train.shape[0], n_folds=folds)
92 for i, (train_index, test_index) in enumerate(kf):
93     print('\n Fold %d\n' % (i + 1))
94     X_train, X_val = train[train_index], train[test_index]
95     y_train, y_val = target[train_index], target[test_index]
96
97 #####
98 #
99 # Define cross-validation variables
100 #
101 #####
102
103     params = {}
104     params['booster'] = 'gbtree'
105     params['objective'] = "reg:linear"
106     params['eval_metric'] = 'mae'
107     params['eta'] = 0.1
108     params['gamma'] = 0.5290

```

[Code](#)[Output \(3\)](#)[Comments \(21\)](#)[Log](#)[Versions \(1\)](#)[Forks \(13\)](#)[Fork Script](#)

```

110     params['subsample'] = 0.9930
111     params['max_depth'] = 7
112     params['max_delta_step'] = 0
113     params['silent'] = 1
114     params['random_state'] = 1001
115
116
117     d_train = xgb.DMatrix(X_train, label=y_train)
118     d_valid = xgb.DMatrix(X_val, label=y_val)
119     watchlist = [(d_train, 'train'), (d_valid, 'eval')]
120
121 #####
122 # Build Model
123 #####
124
125     # Fit model with train / validation

```

```

125         clf = xgb.train(params,
126                         d_train,
127                         100000,
128                         watchlist,
129                         early_stopping_rounds=early_stopping)
130
131 #####
132 # Evaluate Model and Predict
133 #####
134
135     xgb_rounds.append(clf.best_iteration)
136     scores_val = clf.predict(d_valid, ntree_limit=clf.best_ntree_limit)
137     cv_score = mean_absolute_error(np.exp(y_val), np.exp(scores_val))
138     print(' eval-MAE: %.6f' % cv_score)
139     y_pred = np.exp(clf.predict(d_test, ntree_limit=clf.best_ntree_limit))
140
141 #####
142 # Add Predictions and Average Them
143 #####
144
145     if i > 0:
146         fpred = pred + y_pred
147     else:
148         fpred = y_pred
149     pred = fpred
150     cv_sum = cv_sum + cv_score
151
152 mpred = pred / folds
153 score = cv_sum / folds
154 print('\n Average eval-MAE: %.6f' % score)
155 n_rounds = int(np.mean(xgb_rounds))

```

[Code](#)[Output \(3\)](#)[Comments \(21\)](#)[Log](#)[Versions \(1\)](#)[Forks \(13\)](#)[Fork Script](#)

```

158 # Make Full Dataset Predictions
159 #####
160
161 print('\n Training full dataset for %d rounds ...' % n_rounds)
162 watchlist = [(d_train_full, 'train')]
163 clf_full = xgb.train(
164     params, d_train_full,
165     n_rounds,
166     watchlist,
167     verbose_eval=False,)
168 y_pred_full = np.exp(clf_full.predict(d_test))
169
170 # enter the number of iterations from xgb.cv with early_stopping turned on
171 n_fixed = 376

```

1/2

```

173 nfixed = int(n_fixed * (1 + (1. / folds)))
174 print('\n Training full dataset for %d rounds ...\n' % nfixed)
175 clf_fixed = xgb.train(
176     params, d_train_full,
177     nfixed,
178     watchlist,
179     verbose_eval=False,)
180 y_pred_fixed = np.exp(clf_fixed.predict(d_test))
181 timer(start_time)
182
183 print("#\n Writing results")
184 result = pd.DataFrame(mpred, columns=['loss'])
185 result["id"] = ids
186 result = result.set_index("id")
187 print("\n %d-fold average prediction:\n" % folds)
188 print(result.head())
189 result_full = pd.DataFrame(y_pred_full, columns=['loss'])
190 result_full["id"] = ids
191 result_full = result_full.set_index("id")
192 print("\n Full dataset prediction:\n")
193 print(result_full.head())
194 result_fixed = pd.DataFrame(y_pred_fixed, columns=['loss'])
195 result_fixed["id"] = ids
196 result_fixed = result_fixed.set_index("id")
197 print("\n Full dataset (at CV #iterations) prediction:\n")
198 print(result_fixed.head())
199
200 now = datetime.now()
201 score = str(round((cv_sum / folds), 6))
202 sub_file = 'submission_5fold-average-xgb_' + str(score) + '_' + str(

```

[Code](#)[Output \(3\)](#)[Comments \(21\)](#)[Log](#)[Versions \(1\)](#)[Forks \(13\)](#)[Fork Script](#)

```

205 result.to_csv(sub_file, index=True, index_label='id')
206 sub_file = 'submission_full-average-xgb_' + str(now.strftime(
207     "%Y-%m-%d-%H-%M")) + '.csv'
208 print("\n Writing submission: %s" % sub_file)
209 result_full.to_csv(sub_file, index=True, index_label='id')
210 sub_file = 'submission_full-CV-xgb_' + str(now.strftime(
211     "%Y-%m-%d-%H-%M")) + '.csv'
212 print("\n Writing submission: %s" % sub_file)
213 result_fixed.to_csv(sub_file, index=True, index_label='id')

```

[show less](#)

Output

[Download All](#)

3 files

- submission_5fold-ave...
- submission_full-aver...
- submission_full-CV-x...

submission_5fold-average-xgb_1146.10852_2016-10-13-02-40.csv

Download File

id	loss
4	1474.300048828125
6	2071.08837890625
9	7915.3203125
12	5771.5576171875
15	811.8914794921875
17	2246.956298828125
21	2125.173583984375
28	881.2107543945312

Comments

6



Tili · last edited 2 days ago by Tili
2 days ago

This script features **Box-Cox** transformation to correct the skew in continuous variables. XGBoost is used to mop up the mess.

There are 3 different ways of making submission files. One is a simple numeric average of 5 models, each of which is trained on 80% of the data until early stopping criterion is reached. Next, full dataset

Code Output (3) Comments (21) Log Versions (1) Forks (13) Fork Script

Leaderboard:

5fold-average **1118.10987**

full-average 1122.27500

full-CV 1122.42771

At the time of this writing, 5-fold average score is #14 on the leaderboard.

0



modkzs
a day ago

I am little confused why Box-Cox transformation can make such improvement. Can normal distributed feature cause better xgboost performance?

▲
2
▼



LzjtuEr

a day ago

@modkzs, since XGB uses linear regression as the objective, the assumption of linear regression is that the features follow normal distribution.

▲
5
▼



Tili · last edited a day ago by Tili

a day ago

I am little confused why Box-Cox transformation can make such improvement. Can normal distributed feature cause better xgboost performance?

Violations of normal distribution can increase the error in regression analysis. The skew in data is frequently "fixed" by log(N) (or log(N+1)) transformation. However, log(N) transformation is the best only when Box-Cox lambda factor is close to zero. As it is, most of the skewed columns in our dataset have lambda values that are far from zero (see below). There are better **transformations** than log(N) for different values of lambda and Box-Cox procedure finds the best one for a given dataset.

Edit: Below are lambda values of skewed features. log(N) would not be a best transformation for most of them.

-0.585084895197

-0.995197981798

-3.00787936887

Code

Output (3)

Comments (21)

Log

Versions (1)

Forks (13)

Fork Script

-2.2837778627

-2.27901724839

-2.56311666326

-0.613891833259

-0.437372340485

-0.47055182116

-1.44082551851

▲
0
▼

Andrey Shkabko



a day ago

Training full dataset for 366 rounds ...

```
ZeroDivisionError Traceback (most recent call last) in () 162 n_rounds, 163 watchlist, --> 164  
verbose_eval=False,) 165 y_pred_full = np.exp(clf_full.predict(d_test)) 166
```

ZeroDivisionError: integer division or modulo by zero

^^ python 2.7



2

**Tilii**

a day ago

Training full dataset for 366 rounds ...

I have no idea what causes this error. If you look at the **log** file (scroll all the way down), it runs just fine on Kaggle. The same script runs on my computer as well, and I have python 2.7.

Maybe you have one of the packages (or more) that are not recent enough. My setup is:

```
numpy==1.11.0rc2 pandas==0.18.0 scikit-learn==0.17.1 scipy==0.17.0 xgboost==0.4
```

I don't have the latest xgboost, but that shouldn't be a problem as Kaggle does, and this script works on Kaggle.



0

**Andrey Shkabko**

a day ago

```
numpy 1.11.1 pandas 0.18.1 sklearn 0.17.1 scipy 0.18.1 xgboost 0.4
```

Code

Output (3)

Comments (21)

Log

Versions (1)

Forks (13)

Fork Script

17 hours ago

Nice script, I learned a lot. One question, when I train GBMs I usually use subsample of something around .5 because it allows for stable oob error measures and because of this example: http://scikit-learn.org/stable/auto_examples/ensemble/plot_gradient_boosting_regularization.html. Admittedly, I haven't had a lot of experience with them and the example above is just anecdotal. Would you be so kind as to explain the choice of `params['subsample'] = 0.9930`? Thanks again



2

**Tony S**

15 hours ago

I've been thinking about the normality assumption that was brought up here behind the reasoning for the transformation and I think there is something more mysterious going on here.

Isn't the normality assumption in linear regression only assuming that the residual errors are normally distributed? If so, that doesn't necessarily imply that the variables themselves need to be normally

distributed. For example, a simple linear correlation such as $y=2x+\epsilon$ (a small normally distributed error) could have an almost perfect one-to-one correlation and work great in linear regression; x in this case could be any distribution, not just a normal distribution.

I'd love for someone to give me more information or correct me if I am making an incorrect assumption, but it seems to me more likely that you stumbled onto a transformation for one or more of the features that has a better linear correlation with the "loss" variable.

▲
11
▼



Alexandru Papiu

13 hours ago

Nice script and very good score :) I am however a little confused by the comments here. From what I can tell this model is not in any way related to linear regression. It is a boosted model of trees of depth 7. The `['objective'] = 'reg:linear'` line is simply telling xgboost that the target is continuous.

Tree-based models do not make any assumptions on the normality of the features. In fact tree based model are **invariant** to monotonic transformations such as taking logs or box-cox. No matter how crazy the distribution of your feature is, the tree will just pick a point to split on.

Out of curiosity I actually ran the script without any log or box-cox transforms and got very similar results: 1122.44270 with full cv, 1118.89063 with 5-fold average.

I am afraid the good score here comes mostly from ensembling a bunch of well-tuned xgboosts.

▲
5
▼



Faron

12 hours ago

@Alexandru I would change it to: "In fact tree based model are **almost invariant** to monotonic transformations".

[Why monotonic transforms do affect decision trees](#)

▲



DuthonCracker

Code

Output (3)

Comments (21)

Log

Versions (1)

Forks (13)

Fork Script

I have the same question as @AdamWolfeLevin, how did you get the final parameters like `gamma=0.5290`, `min_child_weight=4.2922`?

▲
1
▼



Tili

7 hours ago

explain the choice of `params['subsample'] = 0.9930`? I have the same question as @AdamWolfeLevin, how did you get the final parameters like `gamma=0.5290`, `min_child_weight=4.2922`?

The parameters were found using global Bayesian optimization as implemented [here](#). See [here](#) for an example. Note that most of the code in that example is unnecessary if you run the script locally and have the `bayes_opt` module installed.

1

**Tili**

7 hours ago

I'd love for someone to give me more information or correct me if I am making an incorrect assumption, but it seems to me more likely that you stumbled onto a transformation for one or more of the features that has a better linear correlation with the "loss" variable.

This is quite possible and I think that Alexandru added similar point of view to your argument.

2

**Tili**

7 hours ago

Tree-based models do not make any assumptions on the normality of the features. In fact tree based model are invariant to monotonic transformations such as taking logs or box-cox. No matter how crazy the distribution of your feature is, the tree will just pick a point to split on.

Funny you should mention this, because I was doing Box-Cox feature manipulations as a preparation for linear models you put to such a good use in this [script](#). I tried XGBoost just because it seems to be fashionable these days, and it worked. I am still exploring linear models.

Out of curiosity I actually ran the script without any log or box-cox transforms and got very similar results: 1122.44270 with full cv, 1118.89063 with 5-fold average.

Yet another thing to try would be a gblinear booster in XGBoost.

I am afraid the good score here comes mostly from ensembling a bunch of well-tuned xgboosts.

Though I will have a better feel after completing my linear models, this seems like a valid argument.

Code

Output (3)

Comments (21)

Log

Versions (1)

Forks (13)

Fork Script

0



5 hours ago

@Alexandru I would change it to: "In fact tree based model are almost invariant to monotonic transformations".

But do they always get better? Or they vary for the same reason they vary when the seed is changed?

0

**Tony S**

5 hours ago

Thanks @Alexandru, @Faron, and @Tili for clarification. Makes sense now. The transformation likely has little to no impact on the gbtrees algorithm (or other tree-based algorithms), but it potentially could have some impact on linear models if it provides better linear correlation to loss. I guess I have some food for thought on next steps.

0

**algo.ai**

3 hours ago

@Tilii

Thanks for another nice contribution!

I would like to know roughly `min_child_weight = 4` in XGB is equalivant to what value in sklearn's `GradientBoostingRegressor`'s `min_samples_leaf`?

0

**algo.ai**

3 hours ago

1

**Tilii**

2 hours ago

I would like to know roughly `min_child_weight = 4` in XGB is equalivant to what value in sklearn's `GradientBoostingRegressor`'s `min_samples_leaf`?

I don't know of any simple conversion procedure between `min_child_weight` and `min_samples_leaf`.

It is not that difficult to tune a single parameter, if `min_samples_leaf` is what interests you. See [here](#) and [here](#).

1

**kiddo**

2 hours ago

I can confirm that linear regression does not assume predictors are normally distributed. The Box-Cox transformations are nonlinear and can help with the linearity of the relationship between response and predictor in linear regression, but in a regression tree or RF, there should be no difference except for the seed for RF.

Code

Output (3)

Comments (21)

Log

Versions (1)

Forks (13)

Fork Script

Styling with Markdown supported

Enter your comments.

Post Reply

Code	Output (3)	Comments (21)	Log	Versions (1)	Forks (13)	Fork Script
-------------	------------	---------------	-----	--------------	------------	-------------